



The CMS Data Management System

M. Giffels, Y. Guo, N. Magini,
T. Wildish, V. Kuznetsov

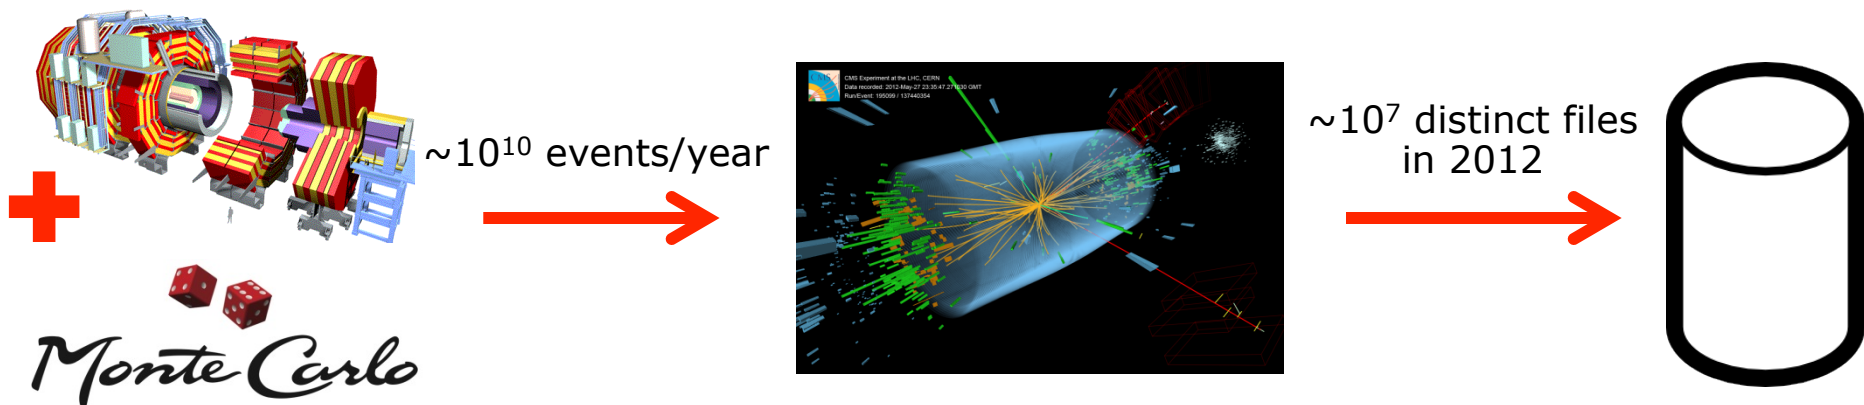


Outline

- Introduction
- Data Management in CMS: performance during LHC Run 1 and planned improvements for LHC Run 2
 - Architecture
 - Core components
 - Data bookkeeping – DBS
 - Data transfers – PhEDEx
 - Query and aggregation service – DAS
 - Additional services
- Summary



CMS Data



- Event data in files
 - average file size reasonably large ~ 2.5 GB
 - Output merged to help scaling in catalogs and storages
- Files are grouped in file blocks to manage them in bulk
 - ~ 10 -1000 files/block
- File blocks are grouped by physics content in datasets of variable size (0.1–100 TB)

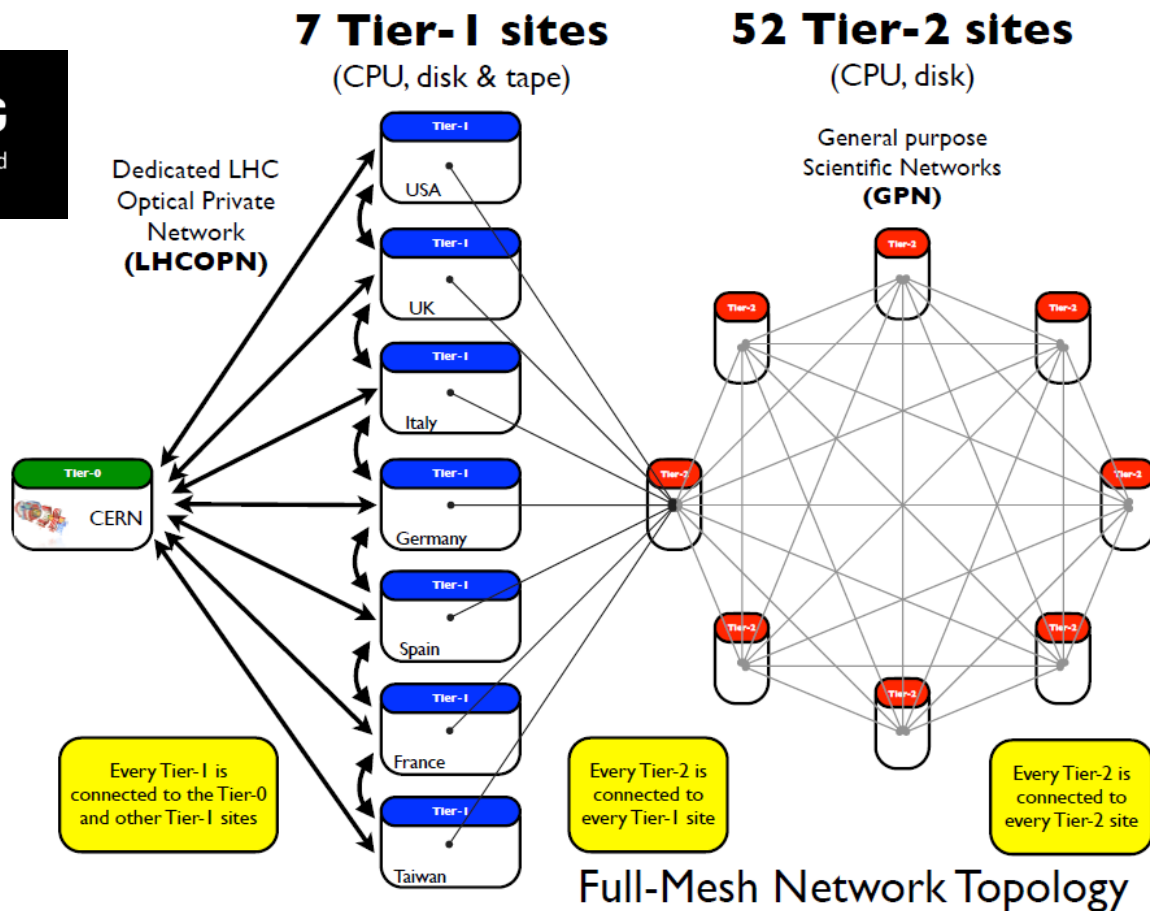


CMS Computing Infrastructure



Tier-0

Custodial storage of first copy of the data
Disk storage for prompt reconstruction and calibration
Distribution of data to Tier-1s



Tier-1s

Shared custodial storage of second copy of the data
Disk storage for data reconstruction and reprocessing
Distribution of analysis data to Tier-2s

Tier-2s

Disk storage for end-user physics analysis and Monte Carlo simulation



Data Management system

- Tasks:
 - Data bookkeeping catalog: *what are the data?*
 - Data location catalog: *where are the data?*
 - Data placement and transfer management
 - For archiving and before processing: CMS job submission is *mostly* based on data location
- CMS Architecture:
 - Independent core components dedicated to different tasks
 - Interacting with each other and with external services via webservices
 - Data from different components aggregated through dedicated service



'Trivial' File Catalog

- *Keep site configuration local*: no global catalog of physical file replicas
- Central services only keep track of the site location of the replicas of logical files
- **TrivialFileCatalog** to map logical to physical file names on local storages: just an xml with a set of regexp rules published by the site
 - Running jobs don't need to contact DM system for file access, they just need to read a local config file
 - Sites can change their storage backend transparently for CMS – just publish a new xml
 - Increased flexibility for sites who need complicated storage setups (e.g. different storages for different namespaces), and great simplicity for sites who don't




Dataset Bookkeeping System (DBS)

- Data bookkeeping: **DBS**
- Data definition
 - What are the data, how they were produced
 - configuration, run number/luminosity, parentage
- Data discovery
 - Which data exist, how they are organized in datasets/blocks/files



Dataset Bookkeeping System (DBS)

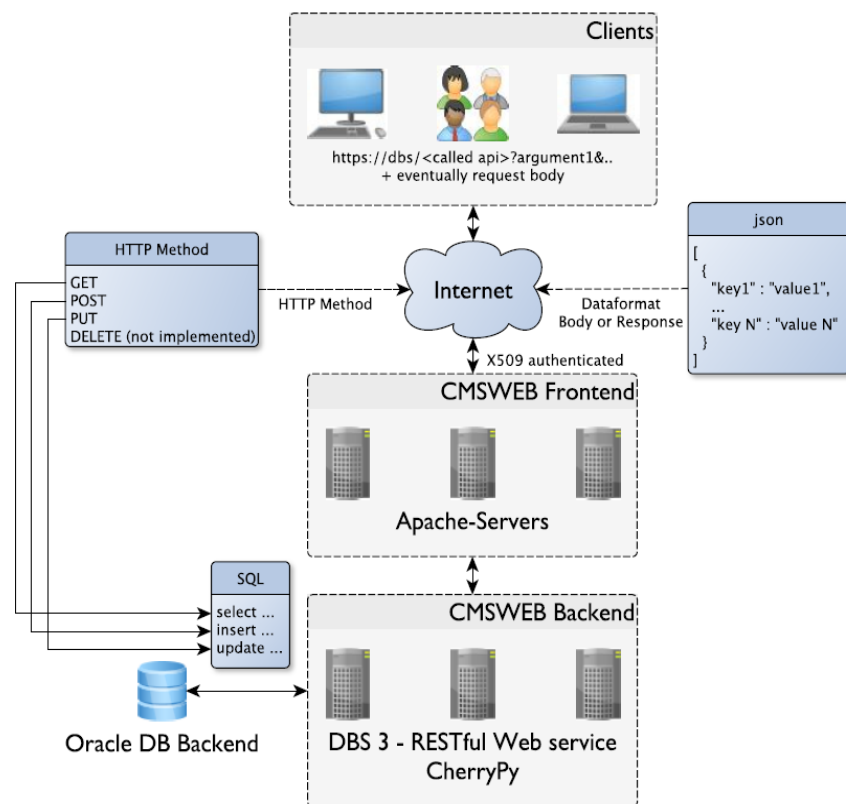
- Current implementation: **DBS2**
 - SQL DB backend
 - Oracle currently used, more supported
 - Multiple instances for different scopes:
Global DBS holds all official CMS data
 - Java servlets under Tomcat
 - XML format for client payload, requires client-side API
- DBS2 sustained the load in LHC Run 1 with some scaling issues
 - At startup the DBS webpage was the main interface for user data discovery, this had to be dropped





Dataset Bookkeeping System (DBS)

- New version: **DBS3**
 - Rescoped schema in Oracle DB to improve scaling, dropping information that did not belong there e.g. data location
 - REST API based on standard cmsweb libs, SQLAlchemy and CherryPy
 - Lightweight information exchange with JSON
- Currently deployed in parallel to DBS2 for validation before final switch¹



1. M. Giffels: Data Bookkeeping Service 3 - Providing event metadata in CMS



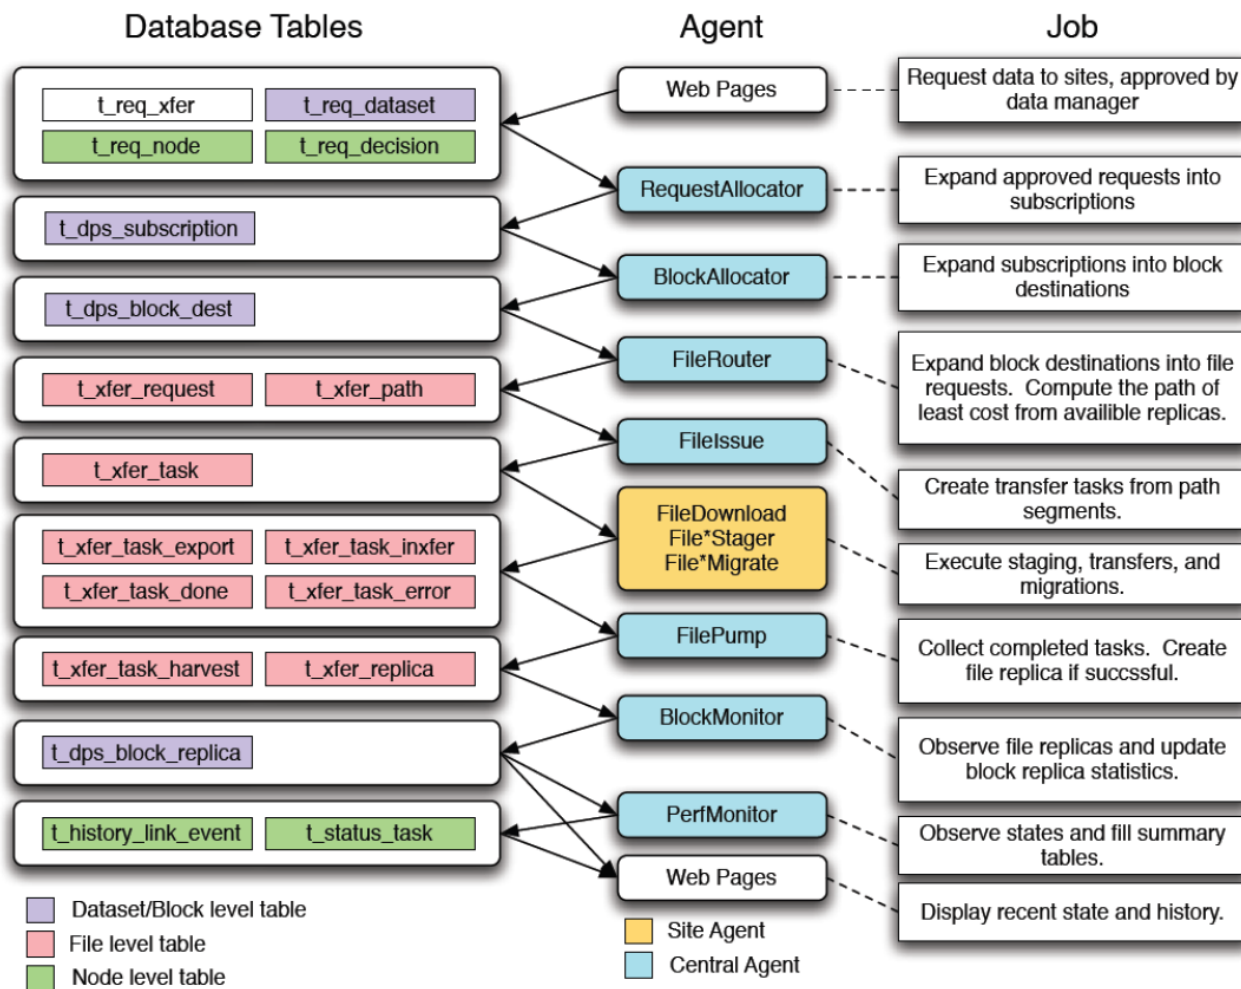
Data location and placement: **PhEDEx** - Physics Experiment Data Export

- Each CMS site runs a set of agents
 - Independent, specialized perl daemons dedicated to fulfill a specific “simple” task
 - Central agents: routing, task assignment, ...
 - Site-specific agents: download, mass storage staging and migration, deletion, consistency checks, ...
- Agents intercommunicate through a central Transfer Management DB (TMDB)
 - Oracle SQL backend
- Web data service and interactive site

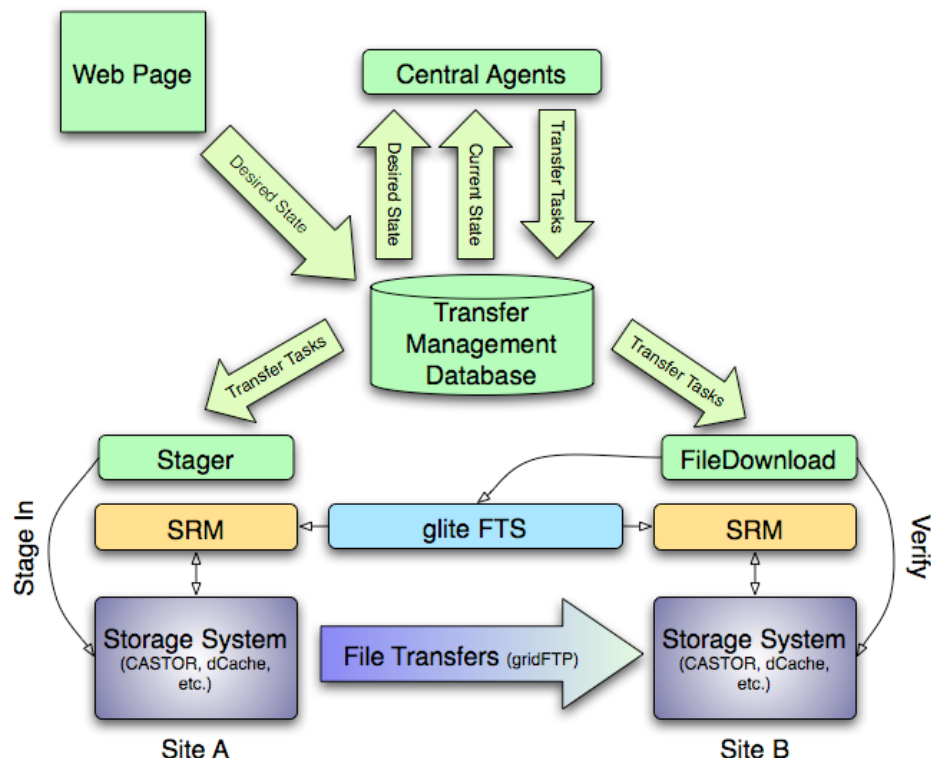
Oracle SQL DB for:

- Replica location catalog
 - Tracked at the level of blocks of files except during transfers, for scaling
 - Only site location of replicas is tracked
- Transfer state blackboard
 - Highly volatile tables for file-level information on active transfer tasks
 - Monitoring tables with aggregated data

PhEDEx workflow




PhEDEx transfer workflow



- Central PhEDEx agents are middleware-agnostic
- Site agents integrated through plugins with WLCG DM middleware – e.g FTS or SRM – to execute transfers



PhEDEx performance



- Up to 77 PB of replicas, ~450k transfers/day in Production



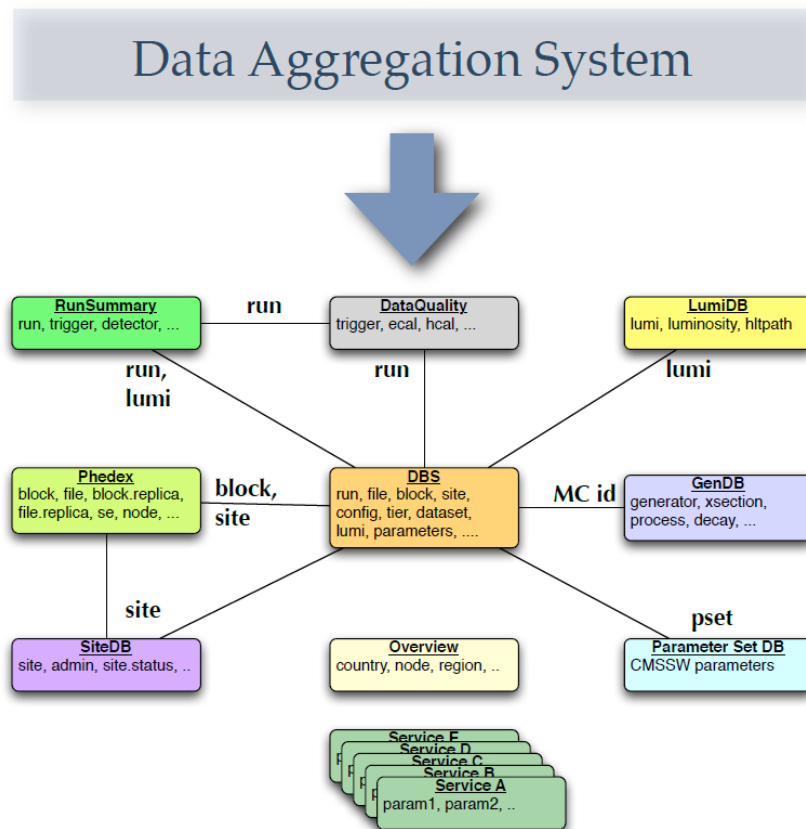
PhEDEx improvements

- PhEDEx scaling proved in simulation testbeds with rates/data volumes $\sim 100x$ higher than production¹
- Recent and planned changes focused on adding support for more flexible workflows
 - Improved support for block transfers
 - Support for requesting generic operator actions e.g. consistency checking²
 - Support for dynamic networking³

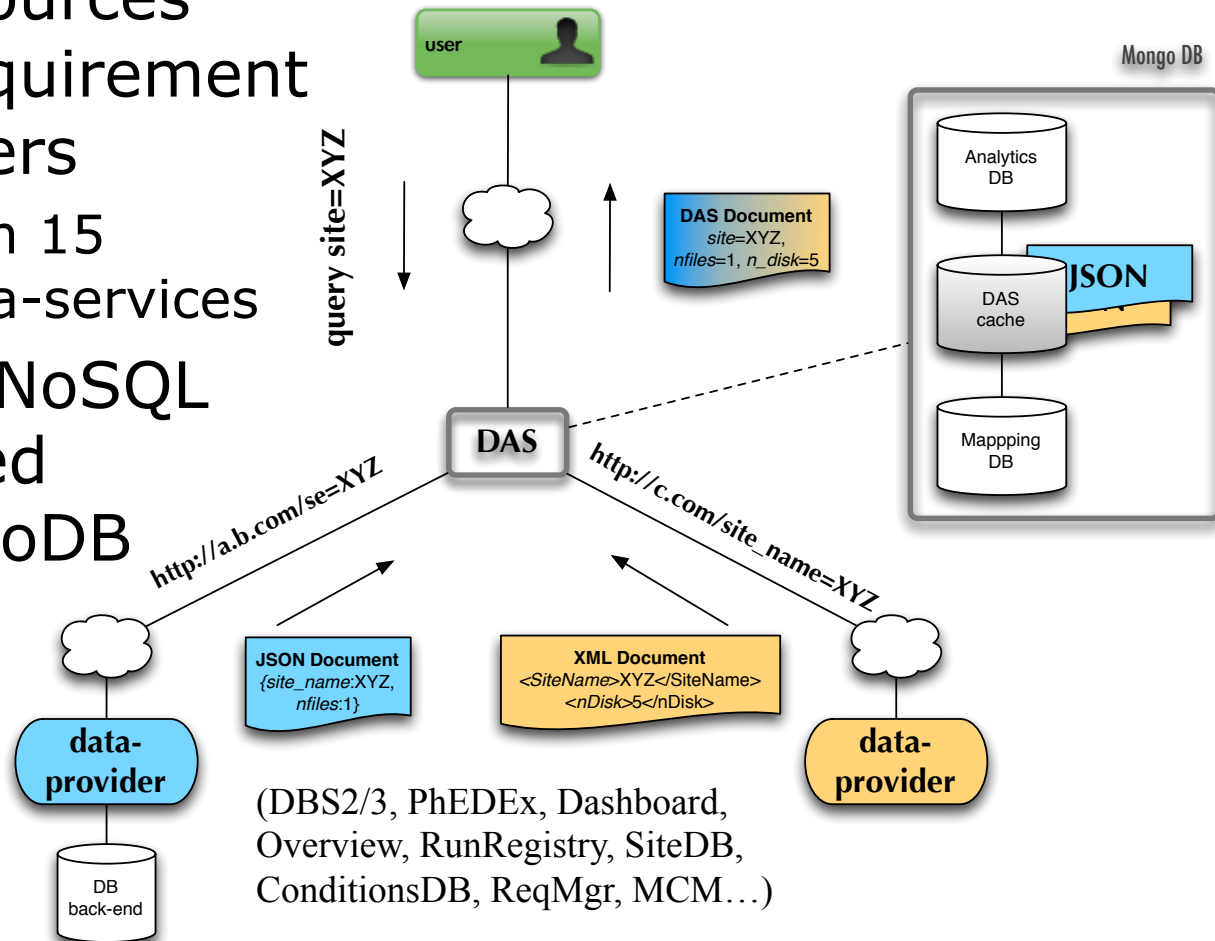
1. T. Wildish: Integration and validation testing for PhEDEx, DBS and DAS with the PhEDEx LifeCycle agent
2. C-H Huang: Request for All - Generalized Request Framework for PhEDEx
3. T. Wildish: Challenging data and workload management in CMS Computing with network-aware systems

Data aggregation

- Deploying independent, dedicated services ensures that each can be optimized for its own task
- Disadvantage: users and services need to query multiple sources to get combined information, e.g.
 - find all sites hosting files for run=XXX
- Solution: **DAS** Data Aggregation System



- Aggregating data from multiple web sources without any requirement on data providers
 - DAS works with 15 distributed data-services
- Data stored in NoSQL document-based database MongoDB
 - For caching and storing aggregated results





DAS interface

The image displays three overlapping screenshots of the Data Aggregation System (DAS) web interface. The top screenshot shows the main navigation bar with links like Home, Services, Bug report, FAQ, CLI, Expert, and Documentation. Below it is a search bar with fields for 'data in', 'format', 'results/per page', and 'records in', along with 'Search' and 'Reset' buttons. The middle screenshot shows a search result for 'site=T1' with a 'Query' section displaying a JSON object for a site record. The bottom screenshot shows an 'Aggregate' section with a 'function: sum' and 'key: site.src_node_files'.

Query

```
{
  "das_id": "4be81e7187eed75f5f000002",
  "id": "4be81f2687eed75f5f000032",
  "site": {
    "kind": "MSS",
    "technology": "dCache",
    "name": "T1_US_FNAL_MSS",
    "se": "cmsrm.fnal.gov",
    "id": "10.0"
  },
  "cust_dest_bytes": "2.71767546431e+15",
  "name": "T1_US_FNAL_MSS",
  "src_node_files": "7924.0",
  "cust_node_bytes": "2.70470672424e+15",
  "cust_dest_files": "1500986.0",
  "nonsrc_node_files": "0.0",
  "nonsrc_node_bytes": "0.0",
  "noncust_dest_files": "867550.0",
  "cust_node_files": "1500682.0",
  "noncust_node_files": "864485.0",
  "src_node_bytes": "9.19710486288e+12",
  "noncust_dest_bytes": "1.83602768692e+15",
  "noncust_node_bytes": "1.82078676773e+15"
}
```

Aggregate

```
{
  "function": "sum",
  "id": 0,
  "result": "1808184.0",
  "key": "site.src_node_files"
}
```


Free text-based query language, filters, aggregators

- DAS interactive webpage fully replaced DBS webpage as main user entry point for data discovery in 2011



DAS performance

- Input queries/day on DAS during 2013
- Each query produces $O(10-1000)$ results
- On average $O(10M)$ results entering and served from DAS cache every day





Additional services

- External services can be easily interfaced to CMS DM components through their web services



CMS POPULARITY & VICTOR

CERN IT Experiment Support

Show 25 entries								Search all columns:	
Cloud	Site	Tier	Group	Total [TB]	Used [TB]	Cleaned [TB]	Full		
AT	Vienna	T2	AnalysisOps	55	81.6	20.9	True		
AT	Vienna	T2	b-tagging	55	55.2	17.7	True		
AT	Vienna	T2	susy	55	51.2	12.8	True		
BE	IIHE	T2	AnalysisOps	274.9	197	None	False		
BE	IIHE	T2	jets-met_hcal	55	45.8	None	False		
BE	IIHE	T2	top	164.9	186.3	46.5	True		
BE	UCL	T2	AnalysisOps	110	142.3	57.6	True		
BE	UCL	T2	exotica	110	161.6	71.1	True		
BE	UCL	T2	tracker-dpg	110	54.4	None	False		
BR	SPRACE	T2	AnalysisOps	219.9	189.5	None	False		

- Victor** data cleaning service deployed in 2011
 - Interfaced with PhEDEx and dataset popularity service to identify unused replicas that can be cleaned up
- Next? Dynamic data placement service to trigger replication in PhEDEx of “hot” (popular) datasets



Summary

- CMS Data Management based on independent core components individually optimized to ensure scaling
- User interface provided by an Aggregation Service that integrates information from the underlying services
- Using a common CMSWEB web service framework allows
 - independent development and evolution of underlying services
 - simplified integration and regression testing when rolling out new service versions
 - building external services that integrate information from several sources in a clean manner