

Integration and validation testing for PhEDEx, DBS and DAS with the PhEDEx LifeCycle agent

C Boeser¹, T Chwalek¹, M Giffels², V Kuznetsov³ and T Wildish⁴

¹ Institut für Experimentelle Kernphysik, Karlsruhe, Germany

² PH-CMG-CO, CERN, CH-1211 Genève 23, Switzerland

³ Cornell University, Ithaca, NY, USA

⁴ Princeton University, Princeton, NJ, USA

E-mail: nicolo.magini@cern.ch

Abstract. The ever-increasing amount of data handled by the CMS dataflow and workflow management tools poses new challenges for cross-validation among different systems within CMS experiment at LHC. To approach this problem we developed an integration test suite based on the LifeCycle agent, a tool originally conceived for stress-testing new releases of PhEDEx, the CMS data-placement tool. The LifeCycle agent provides a framework for customising the test workflow in arbitrary ways, and can scale to levels of activity well beyond those seen in normal running. This means we can run realistic performance tests at scales not likely to be seen by the experiment for some years, or with custom topologies to examine particular situations that may cause concern some time in the future.

The LifeCycle agent has recently been enhanced to become a general purpose integration and validation testing tool for major CMS services (PhEDEx, DBS, DAS). It allows cross-system integration tests of all three components to be performed in controlled environments, without interfering with production services.

In this paper we discuss the design and implementation of the LifeCycle agent. We describe how it is used for small-scale debugging and validation tests, and how we extend that to large-scale tests of whole groups of sub-systems. We show how the LifeCycle agent can emulate the action of operators, physicists, or software agents external to the system under test, and how it can be scaled to large and complex systems.

1. Introduction

PhEDEx [1], DBS [2] and DAS [3] are the core components of the CMS Data Management system [4]. PhEDEx is responsible for moving data around the CMS computing sites, DBS is responsible for maintaining knowledge of the physics content of the data, and DAS is a portal that provides physicists with integrated access to PhEDEx, DBS, and other services.

The development model for these services is highly decoupled. Each is developed by a separate team of people, and each provides a web-based data service for access to information. Ensuring stability and consistency of these interfaces has, until recently, been poorly managed. We have relied on stability of the underlying services and coordination between developers to ensure that successive releases of the software are coherent and correct.

In this paper we describe our solution to this problem. The *LifeCycle agent*, written originally for scale-testing of PhEDEx, has been extended to provide cross-component integrated testing in a flexible and powerful manner. Components can be tested individually or together in different

ways for verification, debugging, or scale and performance testing. The LifeCycle agent itself scales well, and is easily adapted to new requirements (e.g. testing of new APIs from the services). It can inject data into the component systems, retrieve it, manipulate it, and compare it with expectations to detect errors. It can also test expected error conditions for correct behaviour (e.g. access with an unauthorised certificate).

2. The LifeCycle agent

2.1. Basic concepts

The LifeCycle agent is designed to emulate the actions of users or external software processes in a system. The key abstraction is the *workflow*, which consists of a sequence of *events* that depend on each other and are performed at distinct intervals. A sample workflow for PhEDEx, for example, could consist of the following events:

- Generate data - a set of files are generated, organised into file-blocks and one or more datasets
- The files are injected into PhEDEx at a given site - i.e. PhEDEx is told these files exist, and where the initial copy is maintained
- The datasets are subscribed for transfer to one or more other sites
- After a time, the data may be deleted from one or more of the sites

By varying the details of how the generated data is organised, which sites it is sent to, and which sites (if any) it is finally deleted from, we can emulate the pattern of activity in several real-life situations:

- Prompt reconstruction at the Tier-0, from which data is sent to Tier-1s for custodial storage
- Monte Carlo production at Tier-2s, with subsequent upload to a Tier-1, re-distribution to other Tier-2s, and cleanup at the original production site

Many workflows can run in parallel. This means we can emulate the production of dozens of datasets, each destined for different sites, each with different characteristics, to give us a realistic emulation of the total workload of the production transfer system.

2.2. Architecture

2.3. Workflows and Payloads

3. PhEDEx

3.1. Scale tests of PhEDEx

3.2. Testing new PhEDEx releases

The LifeCycle agent is used to test PhEDEx releases. For this purpose a private PhEDEx instance is set up in a testbed with a few nodes, and transfers between these nodes are driven using the LifeCycle agent to orchestrate the flow. In this procedure bug fixes and new features of new PhEDEx versions can be tested before releasing the version.

For special cases, such as bugs that occur under unusual conditions, the LifeCycle agent can be configured to reproduce these conditions, repeatedly and systematically.

3.3. Integration test of the webbased PhEDEx DataService

One of the harder things to test is the web-site access by different users with different access-rights. Rather than requiring the support of several people to systematically test the new release of the website, with all the coordination and overhead that implies, we can use the LifeCycle agent again.

We extend the private testbed to include a website, with self-signed certificates allocated to specific roles. The LifeCycle agent uses the website to perform a series of actions (create

request, approve it, change its priority), repeating this workflow for each certificate, and for the different source and destination nodes. It checks that actions succeed where they should and fail where appropriate. In this way, several combinations of actions (inject, subscribe,...), destination nodes (T0, T1,...), and roles (site-admin, data-manager,...) can be tested in an automated way by configuring and running one single script.

4. DBS

4.1. DBS

blah blah

5. DAS

5.1. Data Aggregation Service

The Data Aggregation Service (DAS) [3] is critical component of CMS Data Management System. It provides ability to query CMS data-services via uniform query language without worrying about security policies and differences in underlying data representations. When physicist poses a question like *"Find me files from RelVal dataset at certain site"*, DAS converts it into series of requests to underlying CMS data-services. Upon processing requested data DAS takes care of aligning naming conventions, data conversion into common JSON format and aggregating data on a requested entity key, e.g. file names. As a result user gets back aggregated record across participated data-providers. This workflow fits really well into LifeCycle design which utilize DAS as cross-check validation tool for its components.

Within LifeCycle Agent DAS was configured to talk to participated DBS and Phedex systems. The queries were processed via DAS engine and presented back to the client. Since original data were requested within LifeCycle Agent itself we were able to compare initial data set with DAS output. This represents a unit test for data injection and retrieval through different systems. We compared the following set of data:

- dataset injection into DBS system
- block and files injection into DBS and Phedex systems
- site content represented in Phedex system

6. Integration testing of PhEDEx, DBS, and DAS

The LifeCycle agent was also used for cross-integration tests of all three components (PhEDEx, DBS3, and DAS) within a controlled environment and thus without interfering with production services. Fake data was created and injected into PhEDEx and DBS. DAS then retrieved information about the data from both sources, and compared the results. For an illustration of the workflow see Fig. 1. By injecting specific errors in either PhEDEx or DBS (changing filenames etc) we can fake errors that we expect to detect with DAS. Special event-handlers are used to compare the errors detected by DAS with the injected errors, and alert us to any unexpected failures.

In the current implementation six kinds of failure can be faked and the probability to occur can be adjusted for each single failure type individually.

- Single file is skipped in the PhEDEx data
- Single file is skipped in the DBS data
- The checksum for a single file is changed for the PhEDEx data
- The checksum for a single file is changed for the DBS data
- The size for a single file is changed for the PhEDEx data
- The size for a single file is changed for the DBS data

The payload-provider (data provider) generates fake data, which then can be injected to a node of the LifeCycle testbed. It also generates extensions to filenames specifying the failure type the file is affected by (according to the probability that was specified).

Special modules read the filenames and act on all files with a certain extension specifying the kind of failure the file is affected by. This includes whether the PhEDEx information, the DBS information, or both are affected and the kind of failure (different checksum, different size or skipped file). Also combinations of these failures are possible for one single file. According to the kind of failure these modules either remove the files, change their size or their checksum, or perform a combination of the mentioned actions, when injecting the information to PhEDEx and DBS3.

After passing these error-handler modules, all files are injected into PhEDEx and DBS and a DAS-query is performed retrieving the information on the injected data from PhEDEx and from DBS3. These two sets of information on the injected data are compared and the found differences are checked against the true information on the injected failures. In this way we can check that DAS properly reports all occurring mismatches in the information on a certain dataset coming both from PhEDEx and/or DBS3.

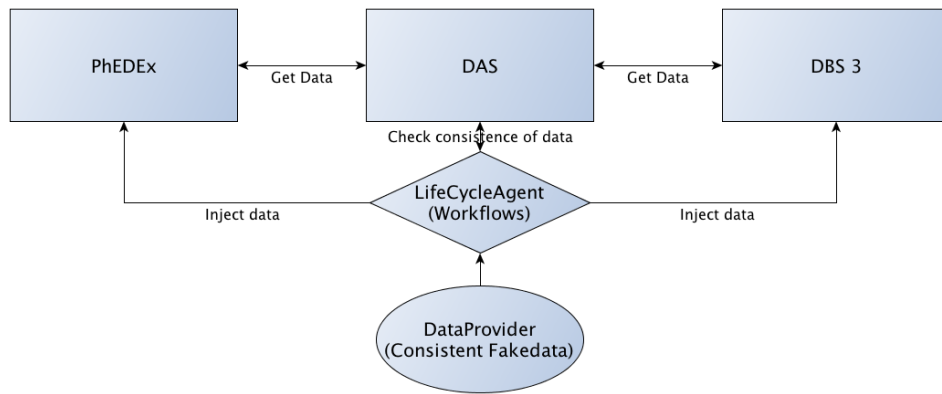


Figure 1. Workflow of the integration testing of PhEDEx, DBS3, and DAS.

7. Conclusions

References

- [1] Egeland R, Wildish T and Metson S 2008 Data transfer infrastructure for CMS data taking, *XII Advanced Computing and Analysis Techniques in Physics Research (Erice, Italy: Proceedings of Science)*
- [2] Giffels M, DBS paper, *submitted to CHEP 2013*
- [3] Kuznetsov V, Evans D and Metson S, The CMS Data Aggregation System, *doi:10.1016/j.procs.2010.04.172*
- [4] Giffels M, Guo Y, Kuznetsov V, Magini N and Wildish T, The CMS Data Management System *submitted to CHEP 2013*